

# Optimizing Thread Pool Performance in Python

Ruotong Chen

Ejiroghene Uwhuba

Xinyuan Jiang

Muhammad Reza

## 1 Introduction

### 1.1 Problem Statement and Importance

Python's `ThreadPoolExecutor` is a widely used tool for managing concurrent tasks, particularly for I/O-bound workloads. However, its performance is often limited by the Global Interpreter Lock (GIL), which prevents multiple native threads from executing Python bytecode in parallel. Understanding the impact of the GIL on threading performance is crucial for optimizing applications that rely on concurrent execution. Additionally, choosing the right thread pool size is a key challenge, as an improperly configured thread pool can lead to resource underutilization or excessive overhead.

Another critical aspect of thread pool performance is task scheduling. The scheduling policy determines how tasks are assigned to worker threads, affecting latency, throughput, and resource efficiency. Common scheduling strategies such as First-In-First-Out (FIFO), Last-In-First-Out (LIFO), priority queues, and round-robin have different trade-offs, but their comparative performance in Python's `ThreadPoolExecutor` has not been extensively studied. Our research aims to evaluate these factors to provide insights into optimizing thread pool performance for different workloads.

### 1.2 Existing Solutions and Their Shortcomings

Several existing studies and best practices discuss when to use `ThreadPoolExecutor` vs. `ProcessPoolExecutor`, highlighting that `ThreadPoolExecutor` is preferable for I/O-bound tasks while `ProcessPoolExecutor` is better for CPU-bound tasks. However, these studies often focus on theoretical guidelines rather than empirical performance benchmarks under real-world workloads.

Regarding thread pool sizing, some general heuristics suggest using formulas like “number of CPU cores  $\times$  2” or “based on the ratio of I/O-bound to CPU-bound tasks,” but these rules do not adapt dynamically to system load, memory constraints, or thread contention. Furthermore, scheduling policies are

often overlooked in Python's standard documentation, despite their potential impact on performance.

### 1.3 Challenges in Addressing This Problem

Optimizing `ThreadPoolExecutor` is challenging because:

- The GIL restricts true parallel execution for CPU-bound tasks, requiring careful benchmarking.
- Optimal thread pool size depends on system resources, workload type, and runtime conditions.
- Different scheduling policies may exhibit varying performance characteristics depending on task duration, arrival patterns, and system contention.
- Python's standard `ThreadPoolExecutor` does not provide built-in scheduling customization, requiring additional implementation effort.

### 1.4 Proposed Solution

Our project will address these issues through a systematic benchmarking study of Python's `ThreadPoolExecutor`. Specifically, we will:

- Analyze the impact of the GIL on threading performance by comparing `ThreadPoolExecutor` with `ProcessPoolExecutor` for different workloads (I/O-bound, CPU-bound, and mixed).
- Investigate thread pool size optimization by testing different pool sizes under various conditions (e.g., fixed size vs. adaptive scaling).
- Evaluate scheduling policies (FIFO, LIFO, priority queue, round-robin) by implementing custom scheduling mechanisms and measuring their impact on task completion time, throughput, and fairness.

## 2 Anticipated Results

We expect to gain insights into:

- The scenarios where `ThreadPoolExecutor` is beneficial or limited due to the GIL.
- The optimal thread pool size based on workload characteristics and system conditions.
- The impact of different scheduling policies on performance, helping developers choose the right strategy for their applications.

Our findings will provide actionable recommendations for improving concurrent performance in Python applications.

## 3 Timeline

Week	Task
Week 1	Literature review on thread pools, GIL, and scheduling policies. Implement baseline benchmarks for <code>ThreadPoolExecutor</code> vs. <code>ProcessPoolExecutor</code> .
Week 2	Experiment with different thread pool sizes and measure performance under varying workloads. Implement custom scheduling mechanisms and compare FIFO, LIFO, priority queues, and round-robin scheduling.
Week 3	Analyze results, refine experiments, and identify key findings. Finalize report, prepare visualizations, and summarize conclusions.

## 4 Evaluation Plan

We will evaluate our solution using:

- **Performance Metrics:**
  - Execution time
  - CPU and memory utilization
  - Task completion latency
  - Throughput (tasks completed per second)
- **Comparative Analysis:**
  - `ThreadPoolExecutor` vs. `ProcessPoolExecutor` for different workloads.
  - Performance variation across different thread pool sizes.
  - Effectiveness of different scheduling policies.